

Having Fun with Gaussian Processes

Emilio Maddalena and Yingzhao Lian

March 2021

1 Setting up the stage

We will use Python 3. The package GPflow 2 will be our easy-to-use toolbox, which implements vanilla GPs, sparse GPs, multi-output GPs, and much more. Take a look at www.gpflow.org and <https://gpflow.readthedocs.io/en/develop/> for more details.

In order to make sure everything is in order, try reconstructing a simple $f(x) = \sin(x)$.

We have 30 samples obtained uniformly from $-\pi \leq x \leq \pi$. Also, the dataset is contaminated with zero-mean Gaussian noise with variance 0.1. Note that the GP doesn't have access to this latter value, and will try to estimate it as a hyperparameter.

Run `simple_test.py` and see how well a squared-exponential GP performs in this task.

Observe how its hyperparameters change after the training phase.

What can we learn from a simple sine wave?

What is the minimum number of samples you can have while keeping a “good fit”?

The range of our ground-truth is $[-1, 1]$. Crank up the noise variance to 1 and look at the output dispersion. Run your code several times and check the results. Increase the number of samples and see if you can still get a smooth sinusoidal fit with low uncertainty.

Let's incorporate some prior knowledge into our experiment. Assume you exactly know the noise level. Use the commands

```
my_gp.likelihood.variance.assign(sigma)
set_trainable(my_gp.likelihood.variance, False)
```

to tell it to the optimizer and to keep it fixed throughout the optimization procedure. Can you now get away with less datapoints?

What if you only had a rough idea of the noise variance, and only wanted to bias the optimization results towards it?

2 A slightly more difficult problem

In `gt_funcs.py` we have defined a number of interesting test functions for you to experiment with. Some are even defined for an arbitrary number of input dimensions.

Run `hard_test.py` and try to reconstruct the Schwefel function in $d = 2$ dimensions

$$f(x) = 418.98d - \sum_{i=1}^d x_i \sin(\sqrt{|x_i|})$$

You will need to pick an appropriate kernel and decide on a minimum number of samples.

3 Playing with sparse models

Run `sparse_test.py` to get a feeling for how easy/difficult it is to train a sparse Gaussian process. You will need to select the number of inducing points M . Both FITC and VFE are implemented in GPflow respectively as GPRFITC and SGPR. If necessary, provide the exact GP hyperparameters as initial guesses to the SGP training phase. For a good comparison between both techniques, we recommend the very good paper (2016) "Understanding probabilistic sparse Gaussian process approximations" by Bauer, Van der Wilk and Rasmussen.

4 Some additional pointers

What is automatic relevance determination (ARD) and in which scenario is it useful?

Take a closer look at the widely used squared-exponential (SE) kernel, otherwise known as the radial basis function one:

$$k(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\ell^2}\right)$$

which belongs to the class of 'stationary' or 'translational-invariant' kernels. Can you imagine a 1D function that would be very hard to describe with a SE-GP? This is perhaps the number one modeling limitation of this class.

Go back to our sine-wave example and replace the input grid by `xx = np.linspace(xmin * 5, xmax * 5, 1000).reshape(-1, 1)`. Run the code again. This suggests that squared-exponential GPs do not extrapolate well. If extrapolation is what you need, you might want to read about semi-parametric Gaussian processes.

Some active researchers in the field: Carl E. Rasmussen, Richard Turner, Mark van der Wilk, Neil D. Lawrence, Robert B. Gramacy, Nando de Freitas, Andrew Gordon Wilson, Andreas Krause, Matthias Seeger, Thomas Schön.